



# Pentesting Cloud Sandboxes in the wild

Matthias Luft & Jan Harrie

# Who we are



**Matthias Luft**  
@uchi\_mata

Principal Platform Security Engineer @Heroku @Salesforce. Opinions are my own. Pentest—Research—Leading—Security Engineering. Love Martial Arts, Outdoors, Dogs.

📅 Seit Februar 2009 bei Twitter

341 Folge ich 771 Follower

Folgen



**jan**  
@NodyTweet

IT-Sec, Container and Orchestration, Freelancer

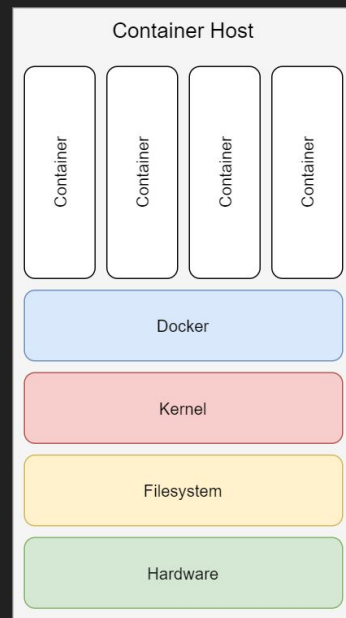
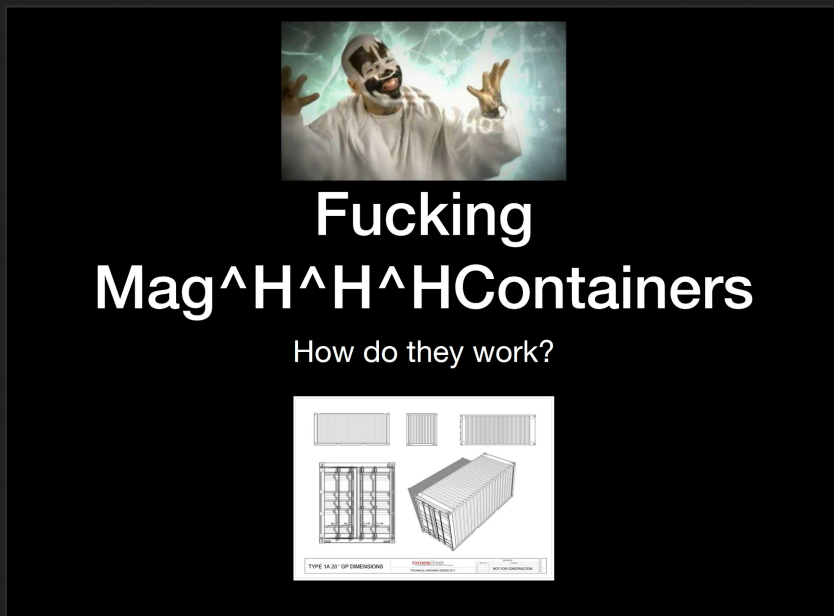
📍 Milliways 🔗 [blog.nody.cc](https://blog.nody.cc) 📅 Seit Dezember 2011 bei Twitter

361 Folge ich 287 Follower

Folgen

# Short Container Re-Cap from last Year

Fucking Containers - how do they work? by Andreas Krebs, BSidesMuc 2019 [1]




# Short Container Re-Cap from last Year

JULIA EVANS  
@b0rk

## what's a container?

a Linux container is a group of processes



Container

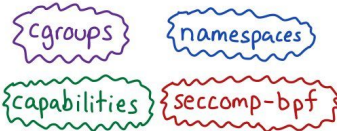
We have our own filesystem but we're still just regular processes!

Linux containers are isolated from other processes

they can have their own:

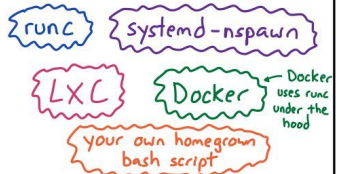
- users
- network namespace
- filesystem
- process IDs
- memory / CPU limits

Kernel features that isolate Linux containers



cgroups namespaces capabilities seccomp-bpf

there are many ways to run Linux containers

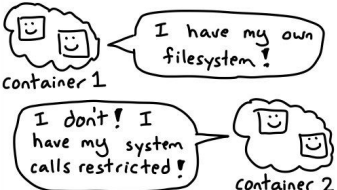


runc systemd-nspawn LXC Docker

your own homegrown bash script

Docker uses runc under the hood

and containers can be set up in different ways



container 1

container 2

**extra confusion:**

"container" sometimes means "lightweight VM"

Fargate and kata Containers are actually VMs and not Linux containers (they don't share a kernel with other containers)

Source: <https://twitter.com/b0rk/status/1225445956734390273>

# What We Will Cover

- Container Breakout Techniques & Attack Surface
- This is a very interesting topic, however:
  - Unless grossly misconfigured, not your most relevant risk.
  - Keep control plane, API, supporting services, and supply chain in mind!

# Attack Vectors – Host Filesystem

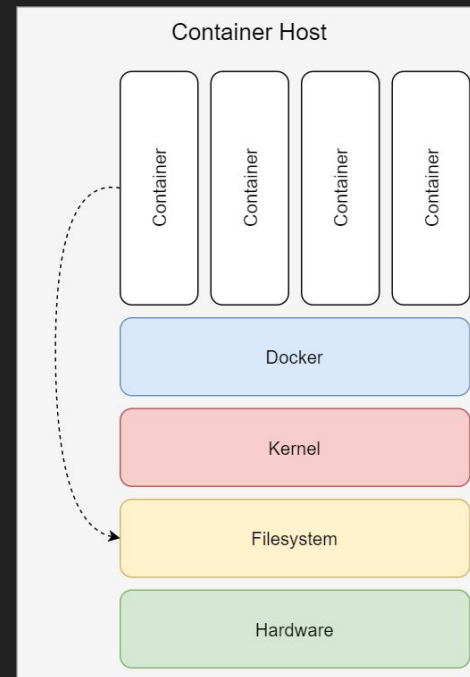
Container is started with access to host root directory

Leverage “normal” Linux privileges escalation techniques

Configure and (ab)use existing host services

⇒ Create user + SSH

⇒ Cronjob



Hands-on blog post “Container Breakouts – Part 1” [2]

# Attack Vectors – Privileged Container

Container is started privileged

Full access to the kernel of the host

⇒ write-access to `/sys`

Leverage Capabilities [7]:

⇒ `CAP_SYS_ADMIN`

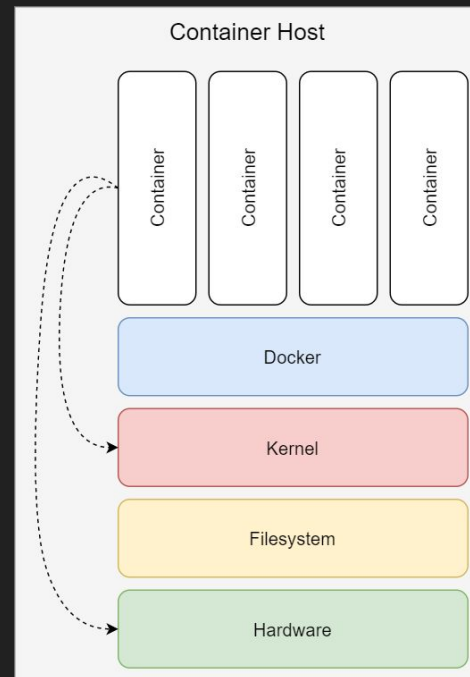
⇒ `CAP_SYS_MODULE`

Host devices

⇒ full access to `/dev`

Host keys [4]

⇒ full access to `/proc/keys`



Hands-on blog post "Container  
Breakouts – Part 2" [3]

# Attack Vectors – Docker Socket

Container is started with mounted Docker Socket (rw)

Full access to container control plain

Access to all running containers

Start “super-power” container

⇒ privileged

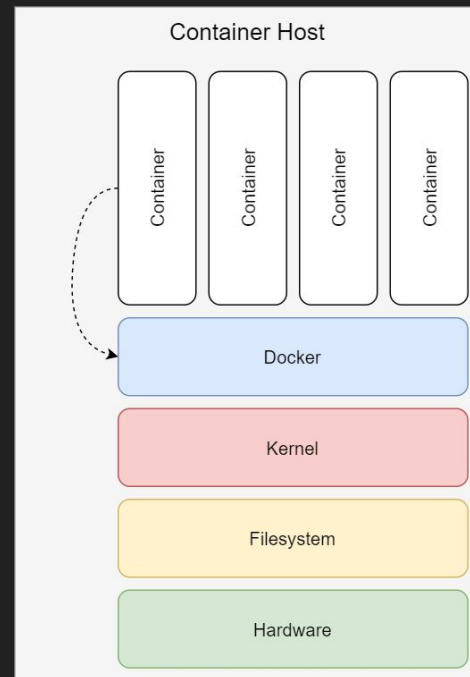
⇒ host network

⇒ host pid/ipc space

⇒ host hostname

⇒ finally enter host filesystem namespace:

`nsenter -t 1 -m`



Hands-on blog post “Container Breakouts – Part 3” [5]



# Attack Vectors – Cloud Interfaces

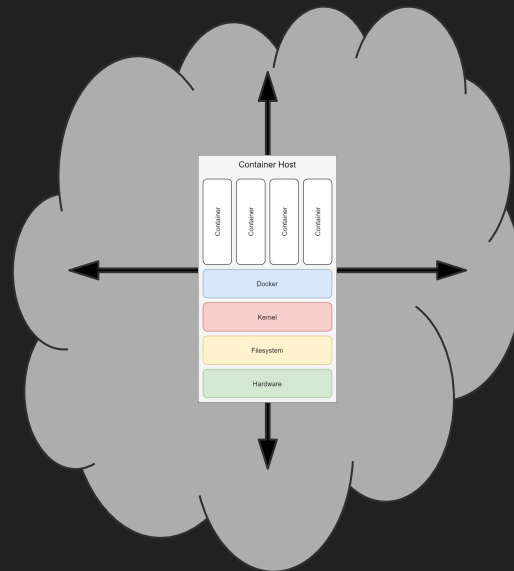
Metadata API available from inside the container

Default configuration generally less impact, **BUT**

*... one is going into the cloud to use cloud features.*

Over permissive configuration leads to.

- Control over workload
- Access to cloud storage
- Cloud account takeover



# Cloud Platform Comparison

Service	PID Namespace	User Namespace	AppArmor Profile/ SELinux	Available Capabilities	Filtered Seccomp	Metadata-Service	Remark
AWS ECS	Yes	No	unconfined	14	0	<a href="http://169.254.169.254/">http://169.254.169.254/</a>	Custom Cluster
AWS Fargate	Yes	No	unconfined	14	65	None	-/-
Azure Container Instances	Yes	No	unconfined	14	65	None	-/-
Docker local	Yes	No	docker-default	13	62	None	Docker Engine
fly.io	No	No	kernel	38	disabled	None	Firecracker
Google Cloud Run	No	No	unconfined	38	disabled	<a href="http://169.254.169.254/">http://169.254.169.254/</a> <a href="http://metadata.google.internal/">http://metadata.google.internal/</a>	gVisor
Heroku	Yes	No	lxc-container-default	23	47	None	LXC

Details can be found: [github.com/NodyHub/bsidesmuc2020](https://github.com/NodyHub/bsidesmuc2020)

# Summary Attack Vectors

- Host Mounts
- Privileges
- Sockets
- Control Plane

# Attack Vectors Automation

- Break-out-the-box, [Botb with a silent T](https://github.com/brompwnie/botb) by the amazing [Chris Le Roy](#)
  - <https://github.com/brompwnie/botb>
- The quite literal `amicontained` by the amazing [Jezz Frazzelle](#)
  - <https://github.com/genuinetools/amicontained/>

# Attack Vectors Automation

- Host Mounts
- Privileges
- **Sockets**
- Control Plane

```
root@7e8f50936949:/# botb -h
[+] Break Out The Box
Usage of botb:
  -aggr string
    Attempt to exploit RuncPWN (default "nil")
  -always-succeed
    Always set BOTB's Exit code to Zero
  -autopwn
    Attempt to autopwn exposed sockets
  -cicd
    Attempt to autopwn but don't drop to TTY,return exit code 1 if successful else 0
  -config string
    Load config from provided yaml file (default "nil")
  -endpoints string
    Provide a textfile with endpoints to use for test (default "nil")
  -find-docker
    Attempt to find Dockerd
  -find-http
    Hunt for Available UNIX Domain Sockets with HTTP
  -find-sockets
    Hunt for Available UNIX Domain Sockets
  -hijack string
    Attempt to hijack binaries on host (default "nil")
  -k8secrets
    Identify and Verify K8's Secrets
  -keyMax int
    Maximum key id range (default 100000000) and max system value is 999999999 (default 100000000)
  -keyMin int
    Minimum key id range (default 1) (default 1)
```

# Attack Vectors Automation

- Host Mounts
- Privileges
- Sockets
- **Control Plane**

```
root@7e8f50936949:/# botb -h
[+] Break Out The Box
Usage of botb:
  -aggr string
    Attempt to exploit RuncPWN (default "nil")
  -always-succeed
    Always set BOTB's Exit code to Zero
  -autopwn
    Attempt to autopwn exposed sockets
  -cicd
    Attempt to autopwn but don't drop to TTY,return exit code 1 if successful else 0
  -config string
    Load config from provided yaml file (default "nil")
  -endpoints string
    Provide a textfile with endpoints to use for test (default "nil")
  -find-docker
    Attempt to find Dockerd
  -find-http
    Hunt for Available UNIX Domain Sockets with HTTP
  -find-sockets
    Hunt for Available UNIX Domain Sockets
  -hijack string
    Attempt to hijack binaries on host (default "nil")
  -k8secrets
    Identify and Verify K8's Secrets
  -keyMax int
    Maximum key id range (default 100000000) and max system value is 999999999 (default 100000000)
  -keyMin int
    Minimum key id range (default 1) (default 1)
```

# Attack Vectors Automation

- Host Mounts
- Privileges
- Sockets
- Control Plane

```
-metadata
  Attempt to find metadata services
-path string
  Path to Start Scanning for UNIX Domain Sockets (default "/")
-pwn-privileged string
  Provide a command payload to try exploit --privilege CGROUP release_agent's (default "nil")
-pwnKeyctl
  Abuse keyctl syscalls and extract data from Linux Kernel keyrings
-recon
  Perform Recon of the Container ENV
-region string
  Provide a AWS Region e.g eu-west-2 (default "nil")
-rev-dns string
  Perform reverse DNS lookups on a subnet. Parameter must be in CIDR notation, e.g., -rev-dns 192.168.0.0/24
(default "nil")
-s3bucket string
  Provide a bucket name for S3 Push (default "nil")
-s3push string
  Push a file to S3 e.g Full command to push to https://YOURBUCKET.s3.eu-west-2.amazonaws.com/FILENAME would
be: -region eu-west-2 -s3bucket YOURBUCKET -s3push FILENAME (default "nil")
-scrape-gcp
  Attempt to scrape the GCP metadata service
-verbose
  Verbose output
-wordlist string
  Provide a wordlist (default "nil")
```

# Attack Vectors Automation

- **Host Mounts**
- **Privileges**
- **Sockets**
- **Control Plane**

```
root@7e8f50936949:/# amicontained
Container Runtime: docker
Has Namespaces:
    pid: true
    user: false
AppArmor Profile: unconfined
Capabilities:
    BOUNDING -> chown dac_override fowner fsetid kill setgid setuid setpca
p net_bind_service net_raw sys_chroot mknod audit_write setfcap
Seccomp: filtering
Blocked Syscalls (64):
    MSGRCV SYSLOG SETSID USELIB USTAT SYSFS VHANGUP PIVOT_ROOT _SYSCTL ACC
T SETTIMEOFDAY MOUNT UMOUNT2 SWAPON SWAPOFF REBOOT SETHOSTNAME SETDOMAINNAME I
OPL IOPERM CREATE_MODULE INIT_MODULE DELETE_MODULE GET_KERNEL_SYMS QUERY_MODUL
E QUOTACTL NFSSERVCTL GETPMSG PUTPMSG AFS_SYSCALL TUXCALL SECURITY LOOKUP_DCOO
KIE CLOCK_SETTIME VSERVER MBIND SET_MEMPOLICY GET_MEMPOLICY KEXEC_LOAD ADD_KEY
REQUEST_KEY KEYCTL MIGRATE_PAGES UNSHARE MOVE_PAGES PERF_EVENT_OPEN FANOTIFY_
INIT NAME_TO_HANDLE_AT OPEN_BY_HANDLE_AT CLOCK_ADJTIME SETNS PROCESS_VM_READV
PROCESS_VM_WRITEV KCMP FINIT_MODULE KEXEC_FILE_LOAD BPF USERFAULTFD MEMBARRIER
PKEY_MPROTECT PKEY_ALLOC PKEY_FREE IO_PGETEVENTS RSEQ
Looking for Docker.sock
```



# Mitigation

- Depends heavily on your runtime
  - Most public PaaS runtime environments won't allow insecure configurations
- If you insist on hosting yourself:
  - Do the same and harden your runtime ;-)
    - Build abstraction layer for deploying workloads with restricted set of options or
    - Lint deployment artefacts

# Mitigation

- Dockerfile:
  - <https://github.com/goodwithtech/dockle>
  - <https://github.com/hadolint/hadolint>
- docker-compose.yml:
  - No tools in our lists, does not seem to be too common anymore.
- Kubernetes YAML:
  - <https://github.com/darkbitio/mkit>
  - <https://www.styra.com>
  - <https://github.com/cruise-automation/k-rail>
  - <https://github.com/zegl/kube-score>
  - Roll your own PSP/OPA/Admission Controllers
    - Good input: <https://aws.github.io/aws-eks-best-practices/pods/>

# Conclusio

- Containers in the cloud can be operated securely
  - ... and are a lot of fun to use from an engineering perspective!
- The challenge is NOT to mis-configure the attack surface
  - *Control plane* also covers provider-specific APIs (see e.g. [6])
- Selection of the cloud platform depends on your flavour
  - ...or on your company policy/golf court ;)

# Further Work

- Cloud Metadata Scraper
- Container privilege footprinting
- Extend botb to scan for binaries with assigned privileges
- Extend amicontained documentation to explain the performed checks

# Q&A

Slides: [gurke.io/bsidesmuc2020](https://gurke.io/bsidesmuc2020)  
Content: [github.com/NodyHub/bsidesmuc2020](https://github.com/NodyHub/bsidesmuc2020)

# References

[1] Fucking Containers - how do they work?, Andreas Krebs, BSidesMuc 2019

[https://2019.bsidesmunich.org/talks/01-03\\_Fucking-Containers/](https://2019.bsidesmunich.org/talks/01-03_Fucking-Containers/)

[https://raw.githubusercontent.com/BSidesMUC/BSidesMunich2019/master/files/01-03\\_Fucking-Containers.pdf](https://raw.githubusercontent.com/BSidesMUC/BSidesMunich2019/master/files/01-03_Fucking-Containers.pdf)

[2] Container Breakouts – Part 1, Jan Harrie

<https://blog.nody.cc/posts/container-breakouts-part1/>

[3] Container Breakouts – Part 2, Jan Harrie

<https://blog.nody.cc/posts/container-breakouts-part2/>

[4] Keyctl-unmask: "Going Florida" on The State Of Containerizing Linux Keyrings, Mark Manning

<https://www.antitree.com/2020/07/keyctl-unmask-going-florida-on-the-state-of-containerizing-linux-keyrings/>

<https://github.com/antitree/keyctl-unmask>

[5] Container Breakouts – Part 3, Jan Harrie

<https://blog.nody.cc/posts/container-breakouts-part3/>

[6] Azure Control Plane Vulnerabilities

<https://research.checkpoint.com/2020/remote-cloud-execution-critical-vulnerabilities-in-azure-cloud-infrastructure-part-i/>

[7] Linux Capabilities - False Boundaries and Arbitrary Code Execution

<https://forums.grsecurity.net/viewtopic.php?t=2522>

# References

- <https://github.com/brompwnie/botb/>
- <https://github.com/guineetools/amicontained/>
- <https://github.com/uchi-mata/system-analysis>